

Building a Prediction Model using ML techniques: GLM,
GAM, Decision Tree, Neural Networks, SVM, Random Forests,
Bagging and Boosting

Professor: Dr. Tony Desmond

Student: Mohsen Selseleh

Student No:

1207477

Fall 2021

https://rpubs.com/Mohsen_Selseleh/848715

Dataset: dvisits (Doctor visits in Australia)

The data come from the Australian Health Survey of 1977-78 and consists of 5190 single adults at 19 variables.

Features	Description
sex	1 if female, 0 if male
age	Age in years divided by 100
agesq	age squared
income	Annual income in Australian dollars divided by 1000
levyplus	1 if covered by private health insurance fund for private patient in public hospital (with doctor of choice), 0 otherwise
freepoor	1 if covered by government because low income, recent immigrant, unemployed, 0 otherwise
freerepa	1 if covered free by government because of old-age or disability pension, or because invalid veteran or family of deceased veteran, 0 otherwise
illness	Number of illnesses in past 2 weeks with 5 or more coded as 5
actdays	Number of days of reduced activity in past two weeks due to illness or injury
hscore	General health questionnaire score using Goldberg's method. High score indicates bad health
chcond1	1 if chronic condition(s) but not limited in activity, 0 otherwise
chcond2	1 if chronic condition(s) and limited in activity, 0 otherwise
doctorco	Number of consultations with a doctor or specialist in the past 2 weeks
nondocco	Number of consultations with non-doctor health professionals (chemist, optician, physiotherapist, social worker, district community nurse, chiropodist or chiropractor) in the past 2 weeks
hospadmi	Number of admissions to a hospital, psychiatric hospital, nursing or convalescent home in the past 12 months (up to 5 or more admissions which is coded as 5)
hospdays	Number of nights in a hospital, etc. during most recent admission: taken, where appropriate, as the mid-point of the intervals 1, 2, 3, 4, 5, 6, 7, 8-14, 15-30, 31-60, 61-79 with 80 or more admissions coded as 80. If no admission in past 12 months then equals zero
medicine	Total number of prescribed and nonprescribed medications used in past 2 days
prescrib	Total number of prescribed medications used in past 2 days
nonpresc	Total number of nonprescribed medications used in past 2 days

Descriptive statistics

sex	age	agesq	income	levyplus	freepoor
Min. :0.0000	Min. :0.1900	Min. :0.0361	Min. :0.0000	Min. :0.0000	Min. :0.000000
1st Qu.:0.0000	1st Qu.:0.2200	1st Qu.:0.0484	1st Qu.:0.2500	1st Qu.:0.0000	1st Qu.:0.000000
Median :1.0000	Median :0.3200	Median :0.1024	Median :0.5500	Median :0.0000	Median :0.000000
Mean :0.5206	Mean :0.4064	Mean :0.2071	Mean :0.5832	Mean :0.4428	Mean :0.04277
3rd Qu.:1.0000	3rd Qu.:0.6200	3rd Qu.:0.3844	3rd Qu.:0.9000	3rd Qu.:1.0000	3rd Qu.:0.000000
Max. :1.0000	Max. :0.7200	Max. :0.5184	Max. :1.5000	Max. :1.0000	Max. :1.000000
freerepa	illness	actdays	hscore	chcond1	chcond2
Min. :0.0000	Min. :0.000	Min. :0.0000	Min. :0.000	Min. :0.0000	Min. :0.0000
1st Qu.:0.0000	1st Qu.:0.000	1st Qu.:0.0000	1st Qu.:0.000	1st Qu.:0.0000	1st Qu.:0.0000
Median :0.0000	Median :1.000	Median :0.0000	Median :0.000	Median :0.0000	Median :0.0000
Mean :0.2102	Mean :1.432	Mean :0.8619	Mean :1.218	Mean :0.4031	Mean :0.1166
3rd Qu.:0.0000	3rd Qu.:2.000	3rd Qu.:0.0000	3rd Qu.:2.000	3rd Qu.:1.0000	3rd Qu.:0.0000
Max. :1.0000	Max. :5.000	Max. :14.0000	Max. :12.000	Max. :1.0000	Max. :1.0000
doctorco	nondocco	hospadmi	hospdays	medicine	prescrib
Min. :0.0000	Min. :0.0000	Min. :0.0000	Min. :0.000	Min. :0.000	Min. :0.0000
1st Qu.:0.0000	1st Qu.:0.0000	1st Qu.:0.0000	1st Qu.:0.000	1st Qu.:0.000	1st Qu.:0.0000
Median :0.0000	Median :0.0000	Median :0.0000	Median :0.000	Median :1.000	Median :0.0000
Mean :0.3017	Mean :0.2146	Mean :0.1736	Mean :1.334	Mean :1.218	Mean :0.8626
3rd Qu.:0.0000	3rd Qu.:0.0000	3rd Qu.:0.0000	3rd Qu.:0.000	3rd Qu.:2.000	3rd Qu.:1.0000
Max. :9.0000	Max. :11.0000	Max. :5.0000	Max. :80.000	Max. :8.000	Max. :8.0000
nonpresc					
Min. :0.0000					
1st Qu.:0.0000					
Median :0.0000					
Mean :0.3557					
3rd Qu.:1.0000					
Max. :8.0000					

Split dataset to train and test datasets

#Set Seed

```
set.seed(10857825)
```

#Training and Testing Data

```
subset2 = sample(nrow(dvisits), nrow(dvisits) * 0.75)
```

```
dvisits.train2 = dvisits[subset2, ]
```

```
dvsits.test2 = dvisits[-subset2, ]
```

#GLM

```
set.seed(10857825)
```

```
glmmodel=glm(doctorco~., data=dvisits.train2)
```

```
summary(glmmodel)
```

Call:

```
glm(formula = doctorco ~ ., data = dvisits.train2)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.7455	-0.2376	-0.1116	-0.0485	7.0186

Coefficients: (1 not defined because of singularities)

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.074615	0.082395	0.906	0.36522

sex	0.017970	0.024753	0.726	0.46790
age	0.143679	0.467794	0.307	0.75875
agesq	-0.172310	0.525018	-0.328	0.74278
income	-0.038144	0.037582	-1.015	0.31019
levyplus	0.006933	0.028176	0.246	0.80566
freepoor	-0.128056	0.059937	-2.137	0.03270 *
freerepa	0.000394	0.042953	0.009	0.99268
illness	0.030791	0.009916	3.105	0.00192 **
actdays	0.090115	0.004275	21.077	< 2e-16 ***
hscore	0.013844	0.005791	2.391	0.01687 *
chcond1	-0.031700	0.027194	-1.166	0.24381
chcond2	-0.058487	0.041421	-1.412	0.15803
nondocco	-0.004674	0.012169	-0.384	0.70097
hospadmi	0.143776	0.026920	5.341	9.78e-08 ***
hospdays	0.002033	0.002219	0.916	0.35967
medicine	-0.027149	0.015925	-1.705	0.08831 .
prescrib	0.120608	0.017754	6.793	1.26e-11 ***
nonpresc	NA	NA	NA	NA

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for gaussian family taken to be 0.4885719)

Null deviance: 2428.3 on 3891 degrees of freedom

Residual deviance: 1892.7 on 3874 degrees of freedom

AIC: 8277.3

Number of Fisher Scoring iterations: 2

extractAIC(glmmodel)

18.000 8277.258

```

extractAIC(glmmodel, k=log(nrow(dvisits.train2)))
18.000 8390.058

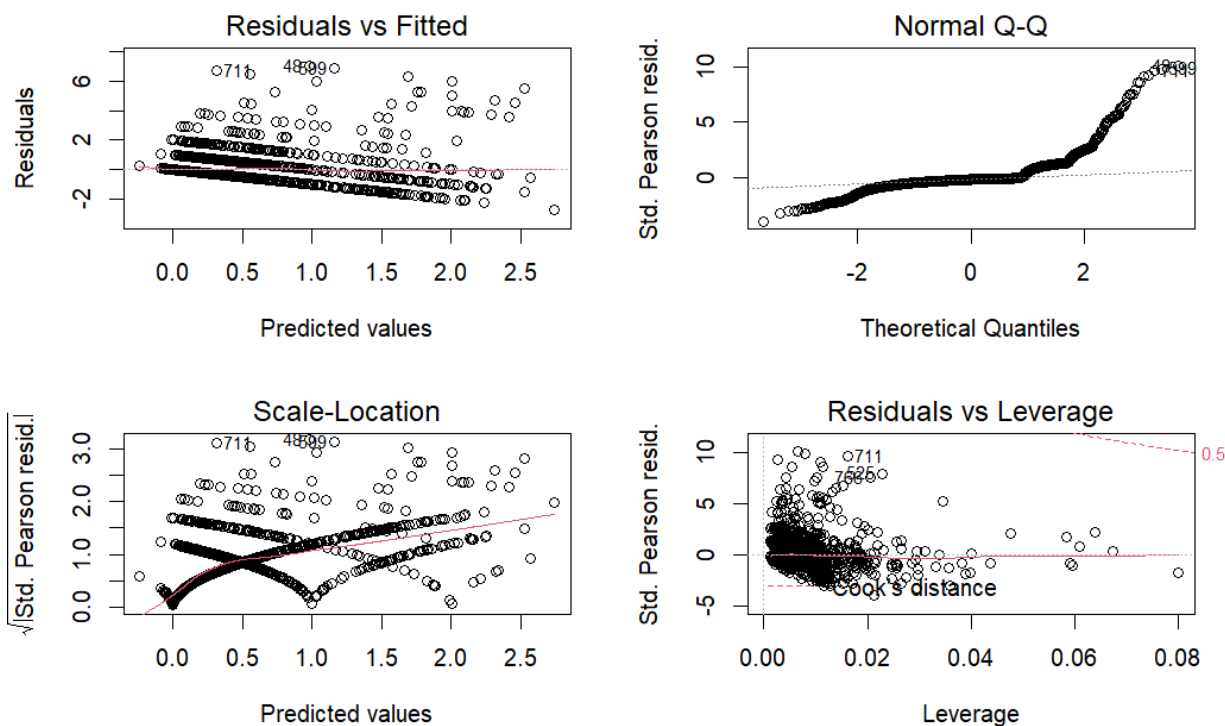
#Prediction with training data
pi1_IS = predict(object = glmmodel, newdata = dvisits.train2)
mean((pi1_IS-dvisits.train2$doctorco)^2)
0.4863123

#Prediction with testing data
pi1_OS = predict(object = glmmodel, newdata = dvisits.test2)
mean((pi1_OS-dvisits.test2$doctorco)^2)
0.4864957

library(modelr)
data.frame(
  R2 = rsquare(glmmodel, data = dvisits),
  RMSE = rmse(glmmodel, data = dvisits),
  MAE = mae(glmmodel, data = dvisits)
)
R2 : 0.2363743
RMSE : 0.6973938
MAE : 0.380663

#Residual Plots
par(mfrow=c(2,2))
plot(glmmodel)

```



cross-validation

set.seed(10857825)

library(boot)

fullmodel = glm(doctorco ~ ., data=dvisits)

cvmodel2<-cv.glm(data = dvisits,glmfit = fullmodel, K = 3)

cvmodel2\$delta[2]

0.4931256

Trees

#Regression Trees

install.packages("rpart")

library(rpart.plot)

library(rpart) #Tree modeling

dvisits.rpart <- rpart(formula = doctorco ~ ., data = dvisits.train2)

```

graph TD
    Root["0.3  
100%"] -- yes --> Node1["0.2  
92%"]
    Root -- no --> Node2["1.4  
8%"]
    Node1 --> Leaf1["0.1  
57%"]
    Node1 --> Leaf2["0.36  
35%"]
    Node2 --> Leaf3["1  
6%"]
    Node2 --> Node3["2.1  
3%"]
    Node3 --> Node4["1.9  
3%"]
    Node3 --> Leaf4["3.9  
0%"]
    Node4 --> Node5["1.3  
1%"]
    Node4 --> Node6["2.3  
2%"]
    Node6 --> Node7["1.6  
1%"]
    Node6 --> Node8["3.1  
1%"]
    Node8 --> Node9["1.3  
0%"]
    Node8 --> Node10["4.3  
0%"]
  
```

```
dvisits.train.pred.tree = predict(dvisits.rpart, dvisits.train2)
mean((dvisits.train.pred.tree - dvisits.train2$doctorco)^2)
0.4417439
```

```
data.frame(
```

```
RMSE = rmse(dvisits.rpart, data = dvisits),
```

R2: 0.2666172

MAE: 0.3811555


```
#Cross-Validation for TREE
```

```
library(caret)
```

```
b <- train(doctorco ~ sex + age + agesq + income + levyplus + freepoor + freerepa  
+ illness + actdays + hscore + chcond1 + chcond2,data=dvisits,
```

```
method = "gam",
```

```
trControl = trainControl(method = "cv", number = 10),
```

```
tuneGrid = data.frame(method = "GCV.Cp", select = FALSE)
```

```
)
```

```
mse<-b$results[3]^2
```

```
mse
```

```
RMSE
```

```
0.5757068
```

```
summary(b$finalModel)
```

```
Family: gaussian
```

```
Link function: identity
```

```
Formula:
```

```
.outcome ~ sex + levyplus + freerepa + chcond1 + chcond2 + illness +  
s(age) + s(agesq) + s(hscore) + s(income)
```

```
Parametric coefficients:
```

```
Estimate Std. Error t value Pr(>|t|)
```

```
(Intercept) 0.121119 0.025639 4.724 2.37e-06 ***
```

```
sex 0.035171 0.023204 1.516 0.130
```

```
levyplus 0.037330 0.026097 1.430 0.153
```

```
freerepa 0.033630 0.040284 0.835 0.404
```

```
chcond1 0.003365 0.025506 0.132 0.895
```

```
chcond2    0.179939  0.038154  4.716 2.47e-06 ***
illness     0.081268  0.008942  9.088 < 2e-16 ***
```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Approximate significance of smooth terms:

	edf	Ref.df	F	p-value
s(age)	1	1	0.013	0.909
s(agesq)	1	1	0.210	0.646
s(hscore)	1	1	74.539	<2e-16 ***
s(income)	1	1	0.710	0.400

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

R-sq.(adj) = 0.0765 Deviance explained = 7.83%

GCV = 0.58952 Scale est. = 0.58827 n = 5190

#GAM

```
require(gam)
```

```
doctor_gam1 <- gam(doctorco ~ sex + age + income + freepoor + freerepa +
illness + actdays + hscore , family=poisson, data = dvisits)
```

```
doctor_gam1
```

```
summary(doctor_gam1)
```

```
AIC(doctor_gam1)
```

```
Call: gam(formula = doctorco ~ sex + age + income + freepoor + freerepa +
illness + actdays + hscore, family = poisson, data = dvisits)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.9423	-0.6818	-0.5757	-0.4942	5.6966

(Dispersion Parameter for poisson family taken to be 1)

Null Deviance: 5634.821 on 5189 degrees of freedom

Residual Deviance: 4388.056 on 5181 degrees of freedom

AIC: 6737.623

Number of Local Scoring Iterations: 5

```
library(modelr)
```

```
data.frame(
```

```
  R2 = rsquare(doctor_gam1, data = dvisits),
```

```
  RMSE = rmse(doctor_gam1, data = dvisits),
```

```
  MAE = mae(doctor_gam1, data = dvisits)
```

```
)
```

R2: 0.06919112

RMSE: 1.936503

MAE: 1.793583

Anova for Parametric Effects

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
sex	1	80.4	80.37	60.8437	7.42e-15 ***
age	1	224.0	223.99	169.5757	< 2.2e-16 ***
income	1	18.0	18.03	13.6492	0.0002226 ***
freepoor	1	8.9	8.88	6.7237	0.0095409 **
freerepa	1	6.9	6.90	5.2232	0.0223274 *
illness	1	469.1	469.11	355.1524	< 2.2e-16 ***
actdays	1	857.1	857.10	648.8860	< 2.2e-16 ***
hscore	1	10.7	10.69	8.0939	0.0044588 **

Residuals 5181 6843.4 1.32

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

[1] 6737.624

BIC(doctor_gam1)

6796.614

#In-sample performance

mean(residuals(doctor_gam1)^2) #In-Sample

0.8454828

pi1 = predict(object = doctor_gam1, newdata = dvisits.test2)

#Out of Sample performance

mean((pi1-dvisits.test2\$doctorco)^2) #Out of Sample

3.804385

#Neural Networks

#Scaling Inputs- To get a range from 0-1

require(neuralnet)

maxs <- apply(dvisits, 2, max)

mins <- apply(dvisits, 2, min)

scaled <- as.data.frame(scale(dvisits, center = mins, scale = maxs - mins))

set.seed(10857825)

subset2 = sample(nrow(dvisits), nrow(dvisits) * 0.75)

dvisits.train2 = scaled[subset2,]

dvisits.test2 = scaled[-subset2,]

n <- names(dvisits.train2)

f <- as.formula(paste("doctorco ~", paste(n[!n %in% "doctorco"], collapse = " +
")))

```

nnetmodel <- neuralnet(f, data=dvisits.train2,hidden=c(8,5), linear.output = TRUE)
plot(nnetmodel)

#Out of Sample

pr.nn<- neuralnet::compute(nnetmodel, dvisits.test2[,1:13])

#Scaling back to get a prediction

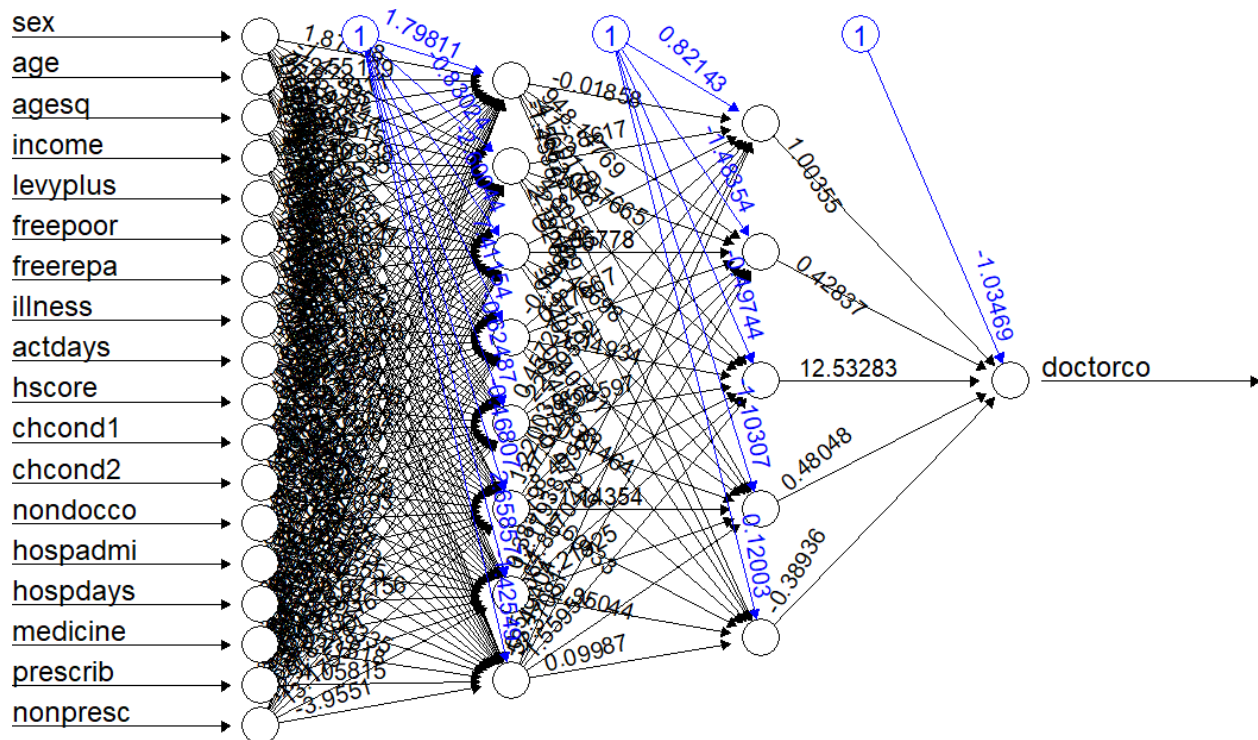
pr.nn_ <- pr.nn$net.result*(max(dvisits$doctorco)-
min(dvisits$doctorco))+min(dvisits$doctorco)

test.r <- (dvisits.test2$doctorco)*(max(dvisits$doctorco)-
min(dvisits$doctorco))+min(dvisits$doctorco)

sum((test.r - pr.nn_)^2)/nrow(dvisits.test2)

...

```



```

## Prediction using neural network

Predict=compute(nnetmodel,dvisits.test2)

mean((Predict$net.result-dvisits.test2$doctorco)^2)

```

0.1624712

```
library(modelr)
```

```
data.frame(
```

```
  R2 = rsquare(nnetmodel, data = dvisits),
```

```
  RMSE = rmse(nnetmodel, data = dvisits),
```

```
  MAE = mae(nnetmodel, data = dvisits)
```

```
)
```

R2: -1.352226

RMSE: 1.227864

MAE: 0.4271425

#Support Vector Machines

```
library(e1071)
```

```
svmmodel<-svm(doctorco~., dvisits.train2)
```

#MSE- In-Sample

```
mean(residuals(svmmodel)^2)
```

```
predsvm<- predict(svmmodel, dvisits.test2)
```

#MSE- Out-of-Sample

```
mean((predsvm-dvisits.test2$doctorco)^2)
```

[1] 0.005431535

[1] 0.007400845

```
library(modelr)
```

```
data.frame(
```

```
  R2 = rsquare(svmmodel, data = dvisits),
```

```
  RMSE = rmse(svmmodel, data = dvisits),
```

```
  MAE = mae(svmmodel, data = dvisits)
```

)

R2: 0.007738408

RMSE: 0.8358721

MAE: 0.3242822

#Bagging

library(randomForest)

set.seed(10857825)

subset2 = sample(nrow(dvisits), nrow(dvisits) * 0.75)

dvisits.train2 = dvisits[subset2,]

dvisits.test2 = dvisits[-subset2,]

set.seed(10857825)

bag.dvisits=randomForest(doctorco~.,data=dvisits.train2,mtry=13,importance=TRUE)

bag.dvisits

Call:

randomForest(formula = doctorco ~ ., data = dvisits.train2, mtry = 13,
importance = TRUE)

Type of random forest: regression

Number of trees: 500

No. of variables tried at each split: 13

Mean of squared residuals: 0.5346554

% Var explained: 14.31

library(modelr)

data.frame(

R2 = rsquare(bag.dvisits, data = dvisits),

```
RMSE = rmse(bag.dvisits, data = dvisits),  
MAE = mae(bag.dvisits, data = dvisits)  
)  
R2: 0.6401244  
RMSE: 0.4790284  
MAE: 0.2515578  
#MSE FOR SVM  
pred.bag = predict(bag.dvisits,newdata=dvisits.test2)  
mean((pred.bag-dvisits.test2$doctorco)^2)  
0.504
```

#Random Forests

```
rf.dvisits=randomForest(doctorco~.,data=dvisits.train2,mtry=6,importance=TRUE)  
rf.dvisits  
Call:  
randomForest(formula = doctorco ~ ., data = dvisits.train2, mtry = 6,  
importance = TRUE)
```

Type of random forest: regression

Number of trees: 500

No. of variables tried at each split: 6

Mean of squared residuals: 0.512795

% Var explained: 17.81

#MSE Random Forest

```
pred.rf = predict(rf.dvisits,newdata=dvisits.test2)  
mean((pred.rf-dvisits.test2$doctorco)^2)  
0.4976903
```




```

library(modelr)
data.frame(
  R2 = rsquare(rf.dvisits, data = dvisits),
  RMSE = rmse(rf.dvisits, data = dvisits),
  MAE = mae(rf.dvisits, data = dvisits)
)
R2: 0.6276496
RMSE: 0.4870887
MAE: 0.2589066

#Boosting
install.packages("gbm")
library(gbm)
set.seed(10857825)
boost.dvisits=gbm(doctorco~.,data=dvisits.train2,distribution="gaussian",n.trees=5000,interaction.depth=4)
summary(boost.dvisits)

```



```
#MSE BOOSTING
```

```
pred.boost=predict(boost.dvisits,newdata=dvisits.test2,n.trees=5000)
```

```
mean((pred.boost-dvisits.test2$doctorco)^2)
```

```
0.6289
```

```
library(modelr)
```

```
data.frame(
```

```
  R2 = rsquare(boost.dvisits, data = dvisits),
```

```
  RMSE = rmse(boost.dvisits, data = dvisits),
```

```
  MAE = mae(boost.dvisits, data = dvisits)
```

```
)
```

```
R2: 0.6136155
```

```
RMSE: 0.4961005
```

```
MAE: 0.2677312
```

```
#GLM WITH POISSON
```

```
doctor_poisson <- glm(doctorco ~ sex + age + agesq + income + levyplus +  
freepoor + freerepa + illness + actdays + hscore + chcond1 + chcond2,  
family=poisson, data = dvisits)
```

```
doctor_poisson
```

```
summary(doctor_poisson)
```

```
par(mfrow=c(2,2))
```

```
plot(doctor_poisson)
```

```
step(doctor_poisson, direction="backward")
```

```
predict(doctor_poisson, dvisits[5190,], type="response")
```

```
print(paste0("Probability of 0 doctor's visits: ", dpois(0, lambda = 0.153)))
```

```
print(paste0("Probability of 1 doctor's visits: ", dpois(1, lambda = 0.153)))
```

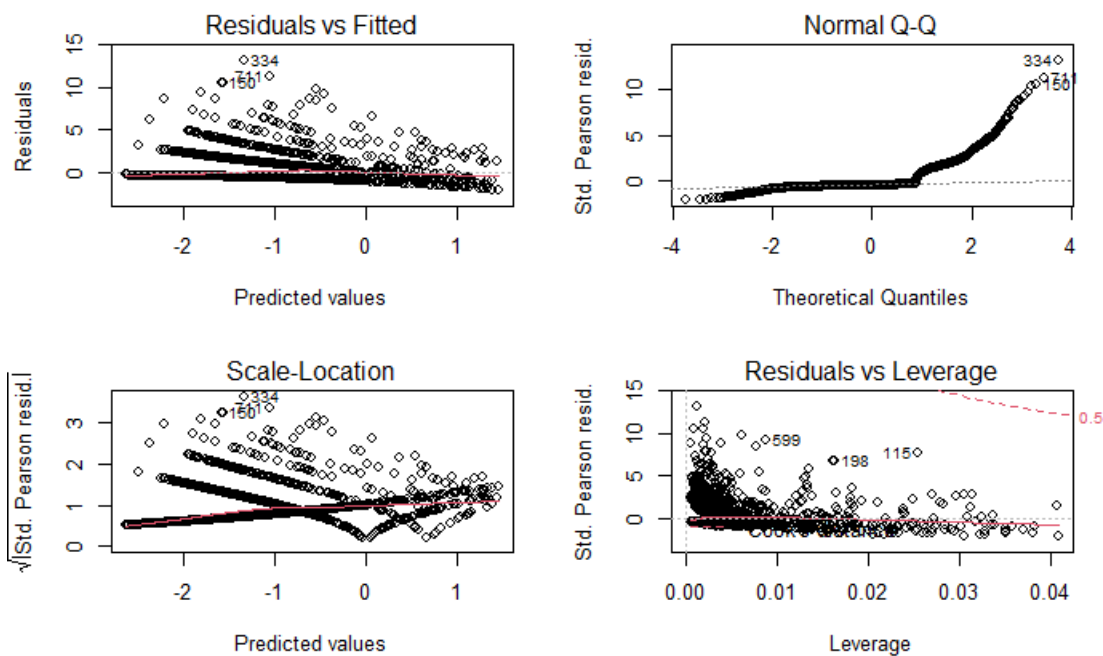
```

print(paste0("Probability of 2 doctor's visits: ", dpois(2, lambda = 0.153)))
print(paste0("Probability of 3 doctor's visits: ", dpois(3, lambda = 0.153)))
print(paste0("Probability of 4 doctor's visits: ", dpois(4, lambda = 0.153)))
print(paste0("Probability of 5 doctor's visits: ", dpois(5, lambda = 0.153)))
round(dpois(0:5, 0.153),3)

doctor_lm <- lm(doctorco ~ sex + age + agesq + income + levyplus + freepoor +
freerepa + illness + actdays + hscore + chcond1 + chcond2, data=dvisits)

doctor_lm
summary(doctor_lm)
predict(doctor_lm, dvisits[5190,])

```



```

# GAM WITH POISSON
#requiring the Package
require(gam)

```

```
doctor_gam <- gam(doctorco ~ sex + age + agesq + income + levyplus + freepoor  
+ freerepa + illness + actdays + hscore + chcond1 + chcond2, family=poisson, data  
= dvisits)
```

```
doctor_gam
```

```
summary(doctor_gam)
```

```
doctor_gam1 <- gam(doctorco ~ sex + age + income + freepoor + freerepa +  
illness + actdays + hscore , family=poisson, data = dvisits)
```

```
doctor_gam1
```

```
summary(doctor_gam1)
```

all:

```
glm(formula = doctorco ~ sex + age + agesq + income + levyplus +  
    freepoor + freerepa + illness + actdays + hscore + chcond1 +  
    chcond2, data = dvisits)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.1352	-0.2588	-0.1435	-0.0433	7.0327

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.027632	0.072220	0.383	0.70202
sex	0.033811	0.021604	1.565	0.11764
age	0.203201	0.410016	0.496	0.62020
agesq	-0.062103	0.458716	-0.135	0.89231
income	-0.057323	0.033089	-1.732	0.08326 .
levyplus	0.035179	0.024882	1.414	0.15748

```

freepoor  -0.103314  0.052471  -1.969  0.04901 *
freerepa   0.033241  0.038157   0.871  0.38371
illness    0.059946  0.008357   7.173  8.39e-13 ***
actdays   0.103192  0.003657  28.216  < 2e-16 ***
hscore     0.016976  0.005190   3.271  0.00108 **
chcond1    0.004384  0.023740   0.185  0.85349
chcond2    0.041617  0.035863   1.160  0.24592

```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for gaussian family taken to be 0.5096272)

Null deviance: 3305.5 on 5189 degrees of freedom

Residual deviance: 2638.3 on 5177 degrees of freedom

AIC: 11245

Number of Fisher Scoring iterations: 2

```
predict(doctor_gam1, dvisits[5190,], type="response")
```

```
print(paste0("Probability of 0 doctor's visits: ", dpois(0, lambda = 0.167)))
```

```
print(paste0("Probability of 1 doctor's visits: ", dpois(1, lambda = 0.167)))
```

```
13.00 11326.58
```

```
print(paste0("Probability of 2 doctor's visits: ", dpois(2, lambda = 0.167)))
```

```
print(paste0("Probability of 3 doctor's visits: ", dpois(3, lambda = 0.167)))
```

```
print(paste0("Probability of 4 doctor's visits: ", dpois(4, lambda = 0.167)))
```

```
print(paste0("Probability of 5 doctor's visits: ", dpois(5, lambda = 0.167)))
```

```
13.00 11326.58
```

```
round(dpois(0:5, 0.153),3)
```

Family: poisson

Link function: log

Formula:

doctorco ~ sex + age + agesq + income + levyplus + freepoor +
freerepa + illness + actdays + hscore + chcond1 + chcond2

Total model degrees of freedom 13

UBRE score: -0.1511532

Family: poisson

Link function: log

Formula:

doctorco ~ sex + age + agesq + income + levyplus + freepoor +
freerepa + illness + actdays + hscore + chcond1 + chcond2

Parametric coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-2.223848	0.189816	-11.716	<2e-16 ***
sex	0.156882	0.056137	2.795	0.0052 **
age	1.056299	1.000780	1.055	0.2912
agesq	-0.848704	1.077784	-0.787	0.4310
income	-0.205321	0.088379	-2.323	0.0202 *
levyplus	0.123185	0.071640	1.720	0.0855 .
freepoor	-0.440061	0.179811	-2.447	0.0144 *
freerepa	0.079798	0.092060	0.867	0.3860
illness	0.186948	0.018281	10.227	<2e-16 ***
actdays	0.126846	0.005034	25.198	<2e-16 ***
hscore	0.030081	0.010099	2.979	0.0029 **
chcond1	0.114085	0.066640	1.712	0.0869 .

chcond2 0.141158 0.083145 1.698 0.0896 .

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

R-sq.(adj) = 0.155 Deviance explained = 22.3%

UBRE = -0.15115 Scale est. = 1 n = 5190

Family: poisson

Link function: log

Formula:

doctorco ~ sex + age + income + freepoor + freerepa + illness +
actdays + hscore

Total model degrees of freedom 9

UBRE score: -0.151049

Family: poisson

Link function: log

Formula:

doctorco ~ sex + age + income + freepoor + freerepa + illness +
actdays + hscore

Parametric coefficients:

Estimate Std. Error z value Pr(>|z|)

(Intercept) -2.052299 0.099492 -20.628 < 2e-16 ***

sex 0.175856 0.055451 3.171 0.00152 **

age 0.448185 0.153571 2.918 0.00352 **

income -0.176135 0.085364 -2.063 0.03908 *

freepoor -0.499016 0.175758 -2.839 0.00452 **

```

freerepa  -0.014633  0.069268  -0.211  0.83270
illness    0.196105  0.017592  11.147  < 2e-16 ***
actdays   0.127813  0.004899  26.087  < 2e-16 ***
hscore     0.032566  0.009961   3.269  0.00108 **

```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

R-sq.(adj) = 0.151 Deviance explained = 22.1%

UBRE = -0.15105 Scale est. = 1 n = 5190

The model GAM is best then the GLM becaues the GAM results hows high probability as compared to the other models

Conclusion

Models	R-Squared	RMSE	MAE
GLM	24	0.69	0.38
GAM	6.9	1.93	1.79
NN	NA	1.23	0.43
SVM	0.0077	0.83	0.32
BAGGING	0.64	0.47	0.25
RF	0.62	0.48	0.26
BOOSTING	0.61	0.5	.27
DT	26.67	.68	.38

Whereas R-squared is a relative measure of fit, RMSE is an absolute measure of fit. Lower values of RMSE indicate better fit. RMSE is a good measure of how accurately the model predicts the response, and it is the most important criterion for fit if the main purpose of the model is prediction. Meanwhile MAE, Lower values are better.

I see that all of them had good performance. But BAGGING, RF, BOOSTING outperformed than other models and they have lower RMSE and MAE. Therefore, I could result that ensemble models were better than other classification models in divisits dataset.